

SE 463 Course Notes

Software Requirements Specification and Analysis

Michael Socha

**4A Software Engineering
University of Waterloo
Spring 2018**

Contents

1	Course Overview	4
1.1	Logistics	4
1.2	Topics Covered	4
2	Business Model	5
2.1	Traditional Product Model	5
2.2	Lean Product Development Model	5
2.3	Lean Canvas	5
2.3.1	Customer Segments	5
2.3.2	Problem Segments	5
2.3.3	Unique Value Proposition	5
2.3.4	Solution	5
2.3.5	Channels	6
2.3.6	Revenue Streams and Costs	6
2.3.7	Key Metrics	6
2.3.8	Unfair Advantage	6
3	Hypothesis Testing	7
3.1	Hypothesis	7
3.2	Falsifiable Hypothesis	7
3.3	Testing Risks	7
3.3.1	Customer Risk	7
3.3.2	Product Risk	7
3.3.3	Market Risk	7
3.4	Problem Interview	8
3.5	Processing Results	8
4	Stakeholders	9
4.1	Owners/Clients/Champions	9
4.2	Customers	9
4.3	Users	9
4.3.1	User Classes	9
4.4	Personas	10
4.5	Domain Expert	10
4.6	Software Engineer	10
4.7	Other Stakeholders	10
5	User Requirements	11
5.1	Use Cases	11
5.2	Actors	11
5.2.1	Actor Generalization	11
5.3	Modifying Use Cases	11
5.4	Use-Case Descriptions	11

5.5	User Stories	11
5.5.1	Benefits of User Stories	12
5.6	Changing Requirements	12
5.6.1	Requirements Baseline	12
5.6.2	Unique Value Proposition (UVP)	12
5.6.3	Project Scope	12
6	Workflow Models	13
6.1	Scenarios	13
6.1.1	More Complex Actions	13
6.2	Workflow Diagrams	14
6.2.1	Activity Diagram	14
6.2.2	Dataflow Diagram (DFD)	14
7	MVP (Minimum Viable Product) Hypothesis Testing	15
7.1	Solution Interview	15
7.1.1	Interview Flow	15
7.1.2	Interview Followup	15
7.2	MVP Advice	16
8	Requirements Elicitation	17
8.1	Artifact-based Elicitation	17
8.1.1	Documents	17
8.1.2	Norms	17
8.1.3	Requirements Taxonomies	17
8.2	Stakeholder-based Elicitation	18
8.2.1	Interviews	18
8.2.2	Questionnaires	18
8.2.3	Ethnographic Analysis	18
8.2.4	Apprenticeship	18
8.2.5	Personas	19
8.3	Model-based Elicitation	19
8.3.1	Requirements Models	19
8.3.2	Mockups and Prototypes	19
8.4	Creativity-based Elicitation	19
8.4.1	Systemic Thinking	19
8.4.2	Brainstorming	20
8.4.3	Creativity Workshops	20
9	Domain Models	21
9.1	Unified Modeling Language (UML)	21
9.2	Conceptual Classes	21
9.3	Finding Domain Entities	21
9.4	Attributes	21
9.5	Associations	22

9.5.1	Roles and Multiplicities	22
9.5.2	Association Classes	22
9.5.3	Association Qualifiers	22
9.6	Composition/Aggregation	22
9.7	Inheritance/Generalization	22
9.8	Common Modeling Conundrums	23
9.8.1	Tightly Coupled Classes	23
9.8.2	Similar but Distinct Classes	23
9.8.3	Representing Object Identity	23
9.8.4	Temporal Objects	23
9.8.5	Class-Scope Attributes	23
10	Prototyping	24
10.1	Prototyping Categories	24
10.2	Prototyping Methodologies	24
10.3	Throwaway Prototyping	24
10.4	Evolutionary Prototyping	24
10.5	Sketching and Wireframing	25
11	Quality Requirements	26
11.1	“Motherhood” Requirements	26
11.2	Fit Criteria	26
11.2.1	Richer Fit Criteria	27
11.3	Measurements	27
11.3.1	Monte Carlo Techniques	27
11.4	Prioritizing Quality Requirements	27

1 Course Overview

1.1 Logistics

- **Professor:** Joanne Atlee

1.2 Topics Covered

Many software project failures are not due to technical reasons, but rather due to poor project specifications. This course focuses on requirement-related topics including:

- Different types of requirements
- Learning about stakeholder needs
- Analyzing and refining preliminary requirements
- Expressing requirements in various representations
- Managing requirement changes

2 Business Model

2.1 Traditional Product Model

A traditional “waterfall” style of development relies heavily on the initial implementation meeting customer needs. This approach often leads to requirement-related problems, since the business and development team might not understand the customer needs or the customer needs might change, and a waterfall model offers little room to refine these requirements.

2.2 Lean Product Development Model

A lean product development model focuses on first determining whether a project addresses an important problem (e.g. problem fit), and whether the project provides a solution people want (market fit). Only after these are established is a product built, which is usually nimble at first and then scales up.

2.3 Lean Canvas

A lean canvas is a template to help determine the initial requirements for a project, especially one in the new product or consulting space. The sections of a lean canvas are described below:

2.3.1 Customer Segments

This is a list of potential customers and users. It helps to have granular segments, and focus should be placed on segments likely to become early adopters.

2.3.2 Problem Segments

This is a list of a few top problems that each customer segment wants solved.

2.3.3 Unique Value Proposition

This is a reason why the project is unique and valuable. The focus should also be on early adopters.

2.3.4 Solution

This is a list of solutions early adopters already apply to solve this problem. From this list, some key capabilities and features of the software project can be determined.

2.3.5 Channels

This is a list of methods to reach the project's target customers. This is important to consider early, since these customers are not just useful as end users, but also as subjects of testing and experiments during development.

2.3.6 Revenue Streams and Costs

Revenue and costs should be considered early in a project. If the project is going to operate on a subscription model, it is a good idea to charge something from the start (with the exception of a free trial period). “Freemium” business models, in which core services are free, tend to have low conversion rates to paying features.

2.3.7 Key Metrics

This is a list of metrics to be used for measuring success.

2.3.8 Unfair Advantage

This refers to some sort of competitive advantage that is hard to replicate. Examples include an existing reputation or customer base, proprietary data, or patented work.

3 Hypothesis Testing

Purpose of hypothesis testing is to validate or disprove guesses in a project's business model. Hypothesis testing is generally conducted by performing pass/fail experiments through market research or customer interviews. Hypothesis being disproved can lead to learning, after which the hypothesis can be pivoted based on new understanding.

3.1 Hypothesis

A hypothesis is a tentative guess about a phenomenon of interest that is testable and falsifiable. The phenomenon of interest (e.g. characteristic of customer problem) are considered to be dependent variables, while independent variables affect the phenomenon of interest. A good hypothesis must be testable (i.e. possible to observe effects of independent variables on dependent variables) and falsifiable.

3.2 Falsifiable Hypothesis

A falsifiable hypothesis is a statement that can easily be proven wrong. Specific and testable statements are more falsifiable than vague ones. Falsifiable hypothesis often have somewhat of a universally quantification aspect (e.g. all, most, etc), rather than merely existential quantification.

3.3 Testing Risks

Uncertainty simply means that there are multiple possibilities. Risk tends to involve uncertainty, with some of the possibilities leading to failure. The three largest risks to a project are:

3.3.1 Customer Risk

e.g. Is there a viable customer segment, and what are their risks?

3.3.2 Product Risk

e.g. How do customers rank the top three problems?

3.3.3 Market Risk

e.g. How do customers solve these problems today?

3.4 Problem Interview

Interviews will generally follow the order below:

- Welcome
- Collect demographics (i.e. test customer segment)
- Tell a story (i.e. set a problem context)
- Problem ranking
- Explore customer's worldview
- Wrapping up (i.e. hook and ask)
- Document results

3.5 Processing Results

It is important to home in on early adopters. As more information is gathered, it may help to refine the problem or pivot the hypothesis. It helps to continue gathering results until they are fairly consistent. Good indications that enough testing has been done are that you can identify the demographics of an early adopter, have a must-have (i.e. critical problem to work on first) problem, and can describe how customers currently solve this problem.

4 Stakeholders

A stakeholder is considered to be anyone who has a stake in a project's ultimate success or failure. Stakeholders include the development team, people in the operational work area, the containing business, and forces from the outside world.

4.1 Owners/Clients/Champions

The owner/client is the group paying for the software to be developed. They are usually considered to be the ultimate/champion stakeholder, since they often have the final say in a project. Examples include a client in a consultancy project, or the company developers are working for in an internal project.

4.2 Customers

A customer buys a project after its completion. The customer may be the same person as the owner or user.

4.3 Users

Users tend to be experts on the work the system is performing, and experts on any existing systems or competing products. Users typically have specific needs a product should satisfy.

4.3.1 User Classes

To meet various user needs, it helps to categorize users by their differences, including:

- Access privilege and security levels
- Tasks regularly performed
- Features used
- Frequency of use
- Application domain expertise, technical expertise
- Platform of use
- Preferred language
- Disfavored users (users who should not have access for security, legal or safety reasons, such as kids on some social media platforms)

4.4 Personas

Personas are resemblances of actual users that can be constructed when real users are hard to interview (e.g. not numerous, too numerous). Personas should imitate the key details of important user classes. Enough details should be provided to make the persona seem realistic. Effectively built personas can:

- Guard against building a product just from the developer’s perspective
- Guard against adapting to “elastic” users, and instead focus on keeners and early adopters

4.5 Domain Expert

Domain experts understand the problem domain very well, and familiar with typical users, their expectations, and potential deployment environments. Note that domain experts need not know much about software engineering.

4.6 Software Engineer

A software engineering is an expert on the project’s development technologies. These stakeholders can represent the development teams they run, and are responsible for overseeing the progress of the technical and some economic aspects of the project. Software engineering can also educate customers how available technology can affect the requested functionality.

4.7 Other Stakeholders

Other potential stakeholders include:

- **Inspectors:** Experts on government and safety regulations
- **Market researchers:** Can serve as proxy for customer interviews
- **Lawyers:** Familiar with legal requirements and licensing
- **Experts on adjacent systems:** Can explain how adjacent systems can affect the project
- **Negative stakeholders:** Stakeholders that want a project to fail

5 User Requirements

5.1 Use Cases

A use case represents some sort of end-to-end functionality in a system, capturing both its triggering event and the complete system response. Well-defined use cases should not overlap much with one another. Time-triggered use cases feature time as their triggering event.

5.2 Actors

An actor is an entity that interacts with the described system, which may include users or other systems. A supporting actor provides some service to the described system.

5.2.1 Actor Generalization

Actors often share common use cases with one another. In these cases, actor generalization can be applied, which factors out common behavior as an abstract actor.

5.3 Modifying Use Cases

- `include` is used to add a sub use case that is used in multiple other use cases
- `extend` is used to add a sub use case that is used in multiple other use cases

Note that well-described systems tend to use these modifiers sparingly.

5.4 Use-Case Descriptions

Textual formats known as use-case descriptions are commonly used to represent use-cases. These descriptions vary in the amount of detail provided for each use case.

5.5 User Stories

User stories are an alternative to use-case descriptions for describing a system. Instead of being told from the system's perspective, user stories provide a description of something a user wants to be able to do (i.e. user requirements). The parts of a successfully managed user story are:

- **Card:** The initial description, often phrased as “As a `role`, I want `something` so that `benefit`”
- **Conversation:** Discussion with product owner to determine requirement details

- **Confirmation:** Criteria for determining whether implementation meets requirements (conditions of satisfaction are described from the system’s perspective)

5.5.1 Benefits of User Stories

The benefits of user stories include:

- Easy for stakeholders to understand
- User-focused style encourages discussion more than ordinary written documentation
- Encouraging iterative development when stories are sized accordingly

5.6 Changing Requirements

System requirements often change during the development process. Below are key topics to consider when managing requirement changes.

5.6.1 Requirements Baseline

A requirements baseline is a set of formally reviewed and agreed core requirements for a system. Major requirement changes should pass through this review process, and it may be helpful to limit the rate of change to these requirements.

5.6.2 Unique Value Proposition (UVP)

New requirements should not weaken a project’s UVP.

5.6.3 Project Scope

When project scope is not managed well, there is a tendency for “scope creep” to occur. Scope creep can be mitigated by keeping a prioritized list of requirements, and only adding a few major ones to each release.

6 Workflow Models

A use case can be decomposed into an ordered sequence of tasks known as a workflow. Each task in a workflow specifies the data it requires and the data it produces.

6.1 Scenarios

A scenario is one full execution path through a use case, listing only observable actions between a system and external actors. A use case can be considered to be a collection of success and failure scenarios initiated by an external actor to achieve a particular goal. An example scenario for a banking system is below:

1. User requests to withdraw funds, specifies amount
2. Bank authenticates user
3. Bank checks account has sufficient funds
4. Bank dispenses cash and receipt
5. Bank records the transaction
6. User takes case and receipt, leaves

A general template for a scenario includes:

- **Business event name**
- **Business use case name/reference**
- **Trigger**
- **Preconditions**
- **Interested stakeholders**
- **Active stakeholders**
- **Normal case steps**
- **Alternatives** - Sub use cases that capture the main goal of this use case through a different sequence of actions
- **Exceptions** - Sub use cases that capture an unwanted deviation
- **Postconditions**

6.1.1 More Complex Actions

Scenario actions can include more complex control flow, such as branching and looping. However, this should be used sparingly, and is often a good indicator that the scenario is

becoming too detailed or all-encompassing.

6.2 Workflow Diagrams

6.2.1 Activity Diagram

Activity diagrams are UML graphical representations of a system's flow control between activities and actions.

6.2.2 Dataflow Diagram (DFD)

DFDs are graphical representations of the functional decomposition of an information system. Each step in a described workflow states what information it requires along with its source, and what information is produced along with its destination.

7 MVP (Minimum Viable Product) Hypothesis Testing

It is important to run experiments to assess the risks of a solution. Areas of risk to assess include customer risk (e.g. who the early adopters are, how they will be reached), and product risk (e.g. whether the solution and UVP align with the customers' main problem, and what the minimum set of features required to launch is).

7.1 Solution Interview

7.1.1 Interview Flow

An interview designed to test the risks of a MVP solution can follow the general flow below:

1. **Welcome:** Set the stage for how the interview works.
2. **Collect demographics:** Try to determine how the interviewee approaches the problem at hand (e.g. for a job search platform, learn about their job search habits).
3. **Tell a story:** Illustrate the top problems the MVP is trying to solve through a brief story.
4. **Demo (Test Solution):** For each problem, go through the workflow and determine whether it addresses the interviewee's main problems. This may appear to be a pitch, but its goal is learning more so than marketing. Attention should be paid to the interviewee's attitude and reactions. Unexpected reactions provide great learning opportunities and should be examined.
5. **Wrapping up:** Ask for permission to follow up, as well as for referrals for other prospective interviewees.

7.1.2 Interview Followup

Based on the results, the earlier pivot hypothesis may need to be pivoted again, after which features may be added/removed/amended. Exit criteria for this series of interviews includes:

- Can identify demographics of early adopters
- Can identify minimum features needed to solve problem
- MVP is significant improvement over existing solutions

7.2 MVP Advice

Before an MVP is built, a UVP should be determined that provides significant value. The MVP should deliver on the UVP, and focus on addressing the customers' number one problem. The MVP should not focus at first on optimizing the solution, but rather on accelerating the learning process to facilitate continual improvement.

8 Requirements Elicitation

Requirements elicitation is the process of collecting requirements for a system. These requirements can be gathered from many types of sources, including stakeholders, existing documentation, similar or previous systems, and the environment of a system. Various elicitation techniques can be applied, including stakeholder analysis, domain analysis, brainstorming, and pilot experimentation.

8.1 Artifact-based Elicitation

The key idea behind artifact-based elicitation is to gather requirements from existing artifacts (e.g. documentation, existing systems) instead of relying on input from stakeholders. Elicitation techniques that are considered artifact-based include document studies, analyzing similar companies, norm studies, domain analysis, and requirements taxonomy.

8.1.1 Documents

Classes of documents that can be used for analysis include:

- **System documentation:** Can include design documents, bug reports, change requests, user manuals, work procedures, usage statistics, marketing data, and performance figures.
- **Environment documentation:** Can include organizational charts, business plans, policy manuals, financial reports, meeting overviews.
- **Domain analysis:** Can include textbooks, surveys, and standards and regulations documentations.

8.1.2 Norms

Requirements can be gathered by forming a comparison to a different existing system (i.e. statements of the form “Build a better X”).

8.1.3 Requirements Taxonomies

Requirements taxonomies are requirements classifications that can act as a list of requirements to be elicited. A sample taxonomy for information systems is an acronym known as PIECES (Performance, Information and data, Economy, Control, Efficiency, and Services). Taxonomies can be classified as being domain-specific (e.g. PIECES is for information systems) or domain-independent.

8.2 Stakeholder-based Elicitation

The key idea behind stakeholder-based elicitation is to gather detailed requirements information specific to particular problems or stakeholders. Elicitation techniques considered to be artifact-based include stakeholder analysis, questionnaires, interviews, observations, task demos, and personas.

8.2.1 Interviews

Interviews can be useful for gathering information specific to particular stakeholders, such as their specific problems or potential areas of conflict. To gather more information, it helps to ask open questions and allowing the stakeholder sufficient time to formulate a response. Examples of open questions include “Why would you use this product?”, or “What problems do you expect this product to solve for you?”. Interviews in groups can help in sharing more ideas, taking the spotlight off specific individuals, and pre-emptively identifying conflicts among stakeholders. Common mistakes during interviews include not interviewing the right people, asking directed questions early on, letting one person dominate a group discussion, or allowing the interview results to over-specify requirements.

8.2.2 Questionnaires

Questionnaires can be useful in gathering insights from a wide group of people, and can consist of open question (to gather suggestions) or closed questions (to gather opinions). A major limitation of questionnaires is that questions need to be fairly concrete, so although gathering large amounts of data can be effective, new insights are hard to obtain.

8.2.3 Ethnographic Analysis

Ethnographic analysis is a study of user behavior through observation, and can be useful in determining the human/social factors in a system that may be hard to gather from interviews alone. Ethnographic analysis can be useful in identifying the critical features of a system, though tends to focus on existing rather than new solutions.

8.2.4 Apprenticeship

An apprentice can gather requirements by spending time with an expert in some field. Information can be gathered through observation, asking questions, and performing some tasks under the expert’s supervision.

8.2.5 Personas

Personas have been previously described; they can be useful if it is not feasible to interview actual users.

8.3 Model-based Elicitation

The key idea behind model-based elicitation is to express requirements languages useful for a particular purpose, which can help raise new insights and concerns. Elicitation techniques considered to be model-based include analysis patterns, mockups and prototyping, and pilot experimentation.

8.3.1 Requirements Models

A model is a simplified representation of a system that captures key details. Models can be used in requirements analysis to guide elicitation, help uncover problems, and provide a measure of progress. Re-expressing requirements through formal models can often reveal problems such as omissions and contradictions.

8.3.2 Mockups and Prototypes

Mockups and prototypes provide the essence of a proposed solution, and are often used to encourage further feedback and requirements elicitation from stakeholders.

8.4 Creativity-based Elicitation

The key idea behind creativity-based elicitation is to come up with new requirements that bring about a competitive advantage through innovation. Elicitation techniques considered to be creativity-based include systemic thinking, brainstorming, creativity workshops, and constraint relaxation.

8.4.1 Systemic Thinking

The goal of systemic thinking is to enlarge the scope of a problem under study from just a system to the outside world at large. The goal of systemic thinking is to gain an understanding of *what* a topic of interest is, with attention paid to its observable behaviors and information flows rather than implementation details.

8.4.2 Brainstorming

Brainstorming is a group activity designed to generate a large number of new and creative ideas. Brainstorming sessions start with an idea generation phase, where participants are encouraged to produce many new ideas, and judgement is suspended. Later, possibly with a new set of participants, the generated ideas are analyzed, and some can be expanded on even further.

8.4.3 Creativity Workshops

Creativity workshops are designed to be spaces for creating new ideas that go above and beyond original stakeholder requirements. Techniques for generating new ideas include brainstorming, constraint relaxation, and analogical reasoning.

9 Domain Models

A domain model is a model of the operating environment of some work under study. Domain models describe real-world entities and their interrelationships, and also capture information about the subject matter and environmental phenomena.

9.1 Unified Modeling Language (UML)

UML is a general-purpose modeling language for software systems. UML provides specifications for various types of structural diagrams (e.g. class diagrams, object diagrams) and behavioral diagrams (e.g. use case diagrams, activity diagrams, state machine diagrams).

UML class diagrams are a good choice for representing domain models. When forming a UML class diagram based on requirements, it is typically best to leave out details; information based on requirements includes class names, key attribute names and association/roles names, while most other information is based on design.

9.2 Conceptual Classes

Conceptual classes provide a high-level overview of what a class does. Conceptual class category examples include business transactions (e.g. FlightReservation), events (e.g. Flight-Departure) and descriptions of things (e.g. FlightDescription).

9.3 Finding Domain Entities

The main ways of finding domain entities include:

- **Refining existing domain models.** Either refining models for previous systems or general domain models.
- **Using a conceptual class list.** Can be generated by brainstorming what kinds of entities and entity classes are relevant to a system.
- **Identifying noun phrases.** Nouns in system descriptions often correspond to entities.

9.4 Attributes

An attribute is a sub-part of an object describing the characteristics or properties of object instances. When forming class diagrams based on requirements, attributes can initially be modeled as just a name, with implementation details added later.

9.5 Associations

Associations represent relevant relationships between multiple conceptual classes. Physical (e.g. “partof”) relationships are typically good candidates, while transient (e.g. “communicates with”) relationships tend to be weaker. Where appropriate, associations can be named, typically with short verbs.

9.5.1 Roles and Multiplicities

A role is one end of a relationship, describing how an entity is viewed within the context of the relationship. Roles can also state multiplicities of association.

9.5.2 Association Classes

Should an association contain attributes of its own, it can be represented with an association class. The state of an association class should only be known by the instances of classes forming the relationship, and each of these objects should have at most one instance of the association class.

9.5.3 Association Qualifiers

Association qualifiers are an association attribute that can be used at one end of an association to distinguish among a set of objects at the other end of the association. An example is a filename serving as an association attribute for a set of files stored in a directory.

9.6 Composition/Aggregation

Composition is a stronger type of association than aggregation. Composition implies that an object cannot exist without its container object (i.e. there is a strong relation to at most one container), while aggregated objects may be weakly related to several containers.

9.7 Inheritance/Generalization

Inheritance is useful for representing “is a kind of” relationships (e.g. an Engineer is a kind of Employee). Properties such as attributes and association of parents can be inherited by children, which can implement abstract features, add new functionality, or override existing functionality. Non-leaf nodes may be abstract classes, which are indicated by italicized names.

9.8 Common Modeling Conundrums

9.8.1 Tightly Coupled Classes

It may be tempting to combine classes with multiple one-to-one composition associations into a single class. However, if the class associations are dynamic, the classes should be kept distinct.

9.8.2 Similar but Distinct Classes

Multiple distinct concepts should not be mixed in a single class. This is particularly important when the concepts have similar names, since this can make them tempting to combine but also produces a confusing model.

9.8.3 Representing Object Identity

Objects should have an identity such that they can be distinguished from other objects with the same values. As such, it is typically a poor idea to base an object's identity on its properties (e.g. identify a house on a street by its street name).

9.8.4 Temporal Objects

In some situations, it is appropriate for a model to store a history of values for an object rather than just the current object.

9.8.5 Class-Scope Attributes

Class-scope attributes have a changeable value that is shared by all object instances of a class. Class-scope attributes can lead to trouble if different instances require that a different value be stored.

10 Prototyping

A prototype is a partial implementation designed to help stakeholders learn more about a problem or its solution.

10.1 Prototyping Categories

Prototypes can be divided into different categories, including:

- **Presentation prototypes**, which are used as a proof of concept to explain certain features/designs and then thrown away.
- **Exploratory prototypes**, which are used to learn more about the problem (e.g. elicit needs, clarify goals), and are also thrown away.
- **Experimental/Breadboard prototypes**, which are used to explore technical feasibility, typically with little to no user/customer involvement.
- **Evolutionary/Operational prototypes**, which is an early deliverable in a continuous development process that gets improved on later.

10.2 Prototyping Methodologies

10.3 Throwaway Prototyping

Throwaway prototypes are used to learn more about a problem, after which they are discarded and not used in a final product. The development of these systems is designed to be done rapidly, with a horizontal strategy often applied that only builds up a single layer at a time (e.g. only build UI for UI-related experimentation). Throwaway prototyping can be useful for gathering information early on (low risk). However, it may result in a lack of formal documentation, unrealistic customer expectations, and in some situations with major time constraints can end up getting developed into a final product (typically leads to poor design).

10.4 Evolutionary Prototyping

Evolutionary prototyping involves learning about a solution to a problem. Emphasis is placed on the “risky” parts of a solution first and the system is developed incrementally, with a vertical strategy that builds up all layers feature-by-feature typically preferred. Evolutionary prototypes tend to be good at evolving to requirements changes, and can be reverted to their previous increment should some change introduce major problems. However, evolutionary prototypes run the risk of poor early architectural choices, poorly structured systems, and often lack high-level direction and control.

10.5 Sketching and Wireframing

Sketches of a UI focus on content, and include the scope of what should be included on a display and how it should be organized. Wireframes are a visual representation of a UI that focus on content hierarchy and flow rather than presentation details (e.g. color, fonts, images, etc.). Sketching and wireframing both allow for the rapid generation of UI designs, which can be iterated on and refined further.

11 Quality Requirements

Functional requirements describe what a system is supposed to do (e.g. what services are provided, I/O behavior). Quality requirements describe some additional constraints for the system.

Common examples of quality requirements include:

- **Performance** (e.g. execution speed, response time, throughput)
- **Reliability** (e.g. fault tolerance, mean time to failure (MTTF))
- **Robustness** (e.g. tolerance to invalid inputs, performance under stress)
- **Adaptability** (e.g. ease of adding new functionality, reusability in other environments)
- **Security** (e.g. controlled access to data, protection against theft)
- **Usability** (e.g. ease of learning, user productivity)
- **Scalability** (performance as workload increases in terms of data size, number of users, etc.)
- **Efficiency** (with respect to utilization of resources)
- **Accuracy/Precision** (e.g. precision of computations, tolerance for computation errors)

11.1 “Motherhood” Requirements

The term motherhood requirements refers to system requirements that are all-encompassing and do not differentiate well between different systems (e.g. “reliable”, “user-friendly”, “maintainable”). Virtually every system must meet such requirements to some degree; what differentiates systems is the degree to which these attributes are required, and their relative importance.

11.2 Fit Criteria

Fit criteria quantify the extend to which a quality requirement must be met. Examples include:

- System must be down no more than 5 minutes a year.
- 90% of users should be able to create an account within 2 minutes.

11.2.1 Richer Fit Criteria

Richer fit criteria are a way of organizing fit criteria that includes a target metric, a minimum metric, and a high-achieving “outstanding” metric. An example is that a system has a minimum response time of 1s, a target response time of 0.5s, and an outstanding response time of 0.1s.

11.3 Measurements

Effectively applying fit criteria can allow for measurement of quality requirements. However, not all fit criteria can be measured on a system prototype or new systems in general (e.g. criteria that concern large amounts of users or long amounts of time). In these cases, possible approaches to measurement include:

- Measuring attributes of prototype.
- Estimating quality attributes.
- Delivering a system and performing more accurate quality measurements later.

11.3.1 Monte Carlo Techniques

Monte Carlo techniques are a class of computational techniques that apply random sampling to generate estimates. Such techniques can be used in estimating various aspects of a system’s requirements. A sample application is deliberately planting N bugs in a system, and then measuring how many are reported by developers or testers. The resulting bugs to reported bugs ratio can be used to give a rough estimate of the total number of unreported bugs in the system.

11.4 Prioritizing Quality Requirements

Different quality requirements can come into conflict with one another. Common areas of conflict include maintainability vs robustness, performance vs security, and performance vs reuse. Conflicts may also occur between quality and functional requirements (e.g. performance vs certain features). In such cases, requirements need to be prioritized.